

CFRL: A Python library for counterfactually fair offline reinforcement learning using data preprocessing

Several Different Contributors^{1¶}

¹ Several Different Departments, University of Michigan ¶ Corresponding author

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) ↗
- [Repository](#) ↗
- [Archive](#) ↗

Editor: [Open Journals](#) ↗

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright
and release the work under a
Creative Commons Attribution 4.0
International License ([CC BY 4.0](#))

Summary

Reinforcement learning (RL) algorithms aim to learn a sequential decision-making rule, often referred to as a “policy”, that maximizes some pre-specified benefit in an environment across multiple or even infinite time steps. It has been widely applied to fields such as healthcare, banking, and autonomous driving. Despite their usefulness, the decisions made by RL algorithms might exhibit systematic bias due to bias in the training data. For example, when using an RL algorithm to assign treatment to patients over time, the algorithm might consistently assign treatment resources to patients of some races at the expense of patients of other races. Concerns have been raised that the deployment of such biased algorithms could exacerbate the discrimination faced by socioeconomically disadvantaged groups.

To address this problem, Wang et al. (2025) extended the concept of single-stage counterfactual fairness (Kusner et al. 2017) to the multi-stage setting and proposed a data preprocessing algorithm that ensures counterfactual fairness in offline reinforcement learning. An RL policy is counterfactually fair if, at every time step, it would assign the same decisions with the same probability for an individual had the individual belong to a different subgroup defined by some sensitive attribute (such as race and gender). At its core, counterfactual fairness views the observed states and rewards as biased proxies of the (unobserved) true underlying states and rewards, where the bias can often be seen as a result of the observed sensitive attribute. In this light, the data preprocessing algorithm ensures counterfactual fairness by removing this bias from the input offline trajectories.

The CFRL library is built upon this definition of RL counterfactual fairness introduced in Wang et al. (2025). It implements the data preprocessing algorithm proposed by Wang et al. (2025) and provides a set of tools to evaluate the value and counterfactual fairness achieved by a given policy. In particular, it takes in an offline RL trajectory and outputs a preprocessed, bias-free trajectory, which could be passed to any off-the-shelf offline RL algorithms to learn a counterfactually fair policy. Additionally, it could also take in an RL policy and return its value and counterfactual fairness metric.

Statement of Need

Many existing Python libraries implement algorithms that ensure fairness in machine learning. For example, Fairlearn and aif360 focus on mitigating bias in single-stage machine learning predictions under statistical association-based fairness criterion such as demographic parity and equal opportunity (). However, they do not accommodate counterfactual fairness, which defines fairness from a causal perspective, and they cannot be easily extended to the reinforcement learning setting in general. Additionally, ml-fairness-gym provides tools to simulate unfairness in sequential decision-making, but it neither implement algorithms that reduce unfairness nor address counterfactual fairness (). To our current knowledge, Wang et al. (2025) is the first

work to study counterfactual fairness in reinforcement learning. Correspondingly, CFRL is also the first code library to address counterfactual fairness in the reinforcement learning setting.

The contribution of CFRL is two-fold. First, it implements a data preprocessing algorithm that removes bias from offline RL training data. For each individual in the data, the preprocessing algorithm estimates the counterfactual states under different sensitive attribute values and concatenates all of the individual's counterfactual states into a new state variable. The preprocessed data can then be directly used by existing RL algorithms for policy learning, and the learned policy should be approximately counterfactually fair. Second, it provides a platform for evaluating RL policies based on counterfactual fairness. After passing in a policy and a trajectory dataset from the target environment, users can assess how well the policy performs in the target environment in terms of the discounted cumulative reward and a counterfactual fairness metric. This not only allows stakeholders to test their fair RL policies before deployment but also offers RL researchers a hands-on tool to evaluate newly developed counterfactually fair RL algorithms.

High-level Design

The CFRL library is composed of 5 major modules. The functionalities of the modules are summarized in the table below.

Module	Functionalities
reader	Implements functions that read tabular trajectory data from either a .csv file or a pandas.DataFrame into a format required by CFRL. Also implements functions that export trajectory data to either a .csv file or a pandas.DataFrame.
preprocessor	Implements the data preprocessing algorithm introduced in Wang et al. (2025).
agents	Implements a fitted Q-iteration (FQI) algorithm, which learns RL policies and makes decisions based on the learned policy. Users can also pass a preprocessor to the FQI; in this case, the FQI will be able to take in unprocessed trajectories, internally preprocess the input trajectories, and directly output counterfactually fair policies.
environment	Implements a synthetic environment that produces synthetic data as well as a simulated environment that simulates the transition dynamics of the environment underlying some real-world RL trajectory data. Also implements functions for sampling trajectories from the synthetic and simulated environments.
fqe	Implements a fitted Q-evaluation (FQE) algorithm, which can be used to evaluate the value of a policy.
evaluation	Implements functions that evaluate the value and counterfactual fairness of a policy. Depending on the user's needs, the evaluation can be done either in a synthetic environment or in a simulated environment.

A general CFRL workflow is as follows: First, simulate a trajectory using environment or read in a trajectory using reader. Then, train a preprocessor using preprocessor to remove the bias in the trajectory data. After that, pass the preprocessed trajectory into the FQI algorithm in agents to learn a counterfactually fair policy. Finally, use functions in evaluation to evaluate the value and counterfactual fairness of the trained policy.

Data Example

We provide a data example to demonstrate how CFRL uses real data to learn a counterfactually fair policy and evaluate the value and counterfactual fairness of the learned policy. This is only one of the many workflows that CFRL can perform. We refer interested readers to the “Example Workflows” section of the CFRL documentation for more workflow examples.

Data Loading

In this demonstration, we use an offline trajectory generated from a `SyntheticEnvironment` using some pre-specified transition rules. Although it is actually synthesized, we treat it as if it is from some unknown environment for pedagogical convenience.

The trajectory contains 500 individuals (i.e. $N = 500$) and 10 transitions (i.e. $T = 10$). The sensitive attribute variable and the state variable are both univariate. The actions are binary (0 or 1) and were sampled using a random policy that selects 0 or 1 randomly with equal probability. The trajectory is stored in a tabular format in a `.csv` file. We load the trajectory from the tabular form into the array format required by CFRL.

```
zs, states, actions, rewards, ids = read_trajectory_from_dataframe(
    path='./data/sample_data_large_uni.csv',
    z_labels=['z1'],
    state_labels=['state1'],
    action_label='action',
    reward_label='reward',
    id_label='ID',
    T=10)
```

We split the trajectory data into a training set (80%) and a testing set (20%). The training set is used to train the policy, while the testing set is used to evaluate the value and counterfactual fairness metric achieved by the policy.

```
(
    zs_train, zs_test,
    states_train, states_test,
    actions_train, actions_test,
    rewards_train, rewards_test
) = train_test_split(zs, states, actions, rewards, test_size=0.2)
```

Preprocessor Training & Trajectory Preprocessing

We now train the preprocessor and preprocess the trajectory. Instead of performing train-test split again, we set `cross_folds=5` to save data and avoid overfitting. In this case, `train_preprocessor()` will internally divide the training data into 5 folds, and each fold is preprocessed using a model that is trained on the other folds. We initialize the `SequentialPreprocessor`, and `train_preprocessor()` will take care of both preprocessor training and trajectory preprocessing.

```
sp = SequentialPreprocessor(z_space=[[0], [1]], num_actions=2, cross_folds=5,
                           mode='single', reg_model='nn')
states_tilde, rewards_tilde = sp.train_preprocessor(
    zs=zs_train, xs=states_train, actions=actions_train, rewards=rewards_train)
```

Policy Learning

Now we train a policy using FQI and the preprocessed data with `sp` as its internal preprocessor. Note that the training data `states_tilde` and `rewards_tilde` are already preprocessed. Thus,

90 we set preprocess=False during training so that the input trajectory will not be preprocessed
91 again by the internal preprocessor (i.e. sp).

```
agent = FQI(num_actions=2, model_type='nn', preprocessor=sp)
agent.train(zs=zs_train, xs=states_tilde, actions=actions_train,
           rewards=rewards_tilde, max_iter=100, preprocess=False)
```

92 SimulatedEnvironment Training

93 Before moving on to the evaluation stage, there is one more thing to do: We need to train a
94 SimulatedEnvironment that mimics the transition rules of the true environment that generated
95 the training trajectory, which will be used by the evaluation functions to simulate the true
96 data-generating environment. To do so, we initialize a SimulatedEnvironment and train it on
97 the whole trajectory data (i.e. training set and testing set combined).

```
env = SimulatedEnvironment(num_actions=2,
                          state_model_type='nn',
                          reward_model_type='nn')
env.fit(zs=zs, states=states, actions=actions, rewards=rewards)
```

98 Value and Counterfactual Fairness Evaluation

99 We now estimate the value and counterfactual fairness achieved by the trained policy
100 when interacting with the target environment using evaluate_value_through_fqe() and
101 evaluate_fairness_through_model(), respectively. The counterfactual fairness is repre-
102 sented by a metric from 0 to 1, with 0 representing perfect fairness and 1 indicating complete
103 unfairness. We use the testing set for evaluation.

```
value = evaluate_reward_through_fqe(zs=zs_test, states=states_test,
                                   actions=actions_test, rewards=rewards_test,
                                   policy=agent, model_type='nn')
cf_metric = evaluate_fairness_through_model(env=env, zs=zs_test, states=states_test,
                                           actions=actions_test, policy=agent)
```

104 The output value is 7.3576775 and cf_metric is 0.041818181818181824, which indicates
105 our policy is close to being perfectly counterfactually fair. Indeed, the CF metric should be
106 exactly 0 if we know the true underlying environment; the reason why it is not exactly 0 here
107 is because we need to estimate the true underlying environment during preprocessing, which
108 can introduce errors.

109 Bonus: Assessing a Fairness-through-unawareness Policy

110 Fairness-through-unawareness proposes to ensure fairness by excluding the sensitive attribute
111 from the agent's decision-making. However, it might still be unfair because of the indirect
112 bias in the states and rewards. In this section, we use the same trajectory data train a policy
113 following fairness-through-unawareness and estimate its value and counterfactual fairness.

```
agent_unaware = FQI(num_actions=2, model_type='nn', preprocessor=None)
agent_unaware.train(zs=zs_train, xs=states_train, actions=actions_train,
                   rewards=rewards_train, max_iter=100, preprocess=False)
value_unaware = evaluate_reward_through_fqe(
    zs=zs_test, states=states_test, actions=actions_test,
    rewards=rewards_test, policy=agent_unaware, model_type='nn')
cf_metric_unaware = evaluate_fairness_through_model(
    env=env, zs=zs_test, states=states_test,
    actions=actions_test, policy=agent_unaware)
```

114 The output value is 8.588442 and cf_metric is 0.44636363636363635. The fairness-through-
115 unawareness policy is much less fair than the policy learned using the preprocessed trajectory.

116 This suggests that the preprocessing method likely reduced the bias in the training trajectory
117 effectively.

118 Conclusions

119 CFRL is a Python library that empowers counterfactually fair reinforcement learning using
120 data preprocessing. It also provides tools to evaluate the value and counterfactual fairness
121 of a given policy. As far as we know, it is the first library to address counterfactual fairness
122 problems in the context of reinforcement learning. Nevertheless, despite this, CFRL also admits
123 a few limitations. For example, the current CFRL implementation requires every episode in the
124 offline dataset to have the same number of time steps. Extending the library to accommodate
125 variable-length episodes can improve its flexibility and usefulness. Besides, CFRL could also be
126 made more well-rounded by integrating the preprocessor with established offline RL algorithm
127 libraries such as d3r1py, or connecting the evaluation functions with popular RL environment
128 libraries such as gym. We leave these extensions to future updates.

129 Acknowledgements

130 This is the acknowledgements.