

CFRL: A Python library for counterfactually fair offline reinforcement learning via sequential data preprocessing

Jianhan Zhang¹, Jitao Wang², Chengchun Shi³, John D. Piette⁴, Donglin Zeng², and Zhenke Wu^{2¶}

¹ Department of Statistics, University of Michigan, USA ² Department of Biostatistics, University of Michigan, USA ³ Department of Statistics, London School of Economics, UK ⁴ Department of Health Behavior and Health Equity, School of Public Health, University of Michigan, USA ¶ Corresponding author

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) ↗
- [Repository](#) ↗
- [Archive](#) ↗

Editor: [Open Journals](#) ↗

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

Reinforcement learning (RL) aims to learn and evaluate a sequential decision-making rule, often referred to as a “policy”, that maximizes expected discounted cumulative rewards to optimize population-level benefit in an environment across possibly infinitely many time steps. RL has gained popularity in fields such as healthcare, banking, autonomous driving, and, more recently, large language model pre-training. However, the sequential decisions made by an RL algorithm may disadvantage individuals with certain values of a sensitive attribute (e.g., race/ethnicity, income, gender, education level). An RL algorithm learns an optimal policy that makes decisions based on observed state variables. If certain values of the sensitive attribute influence the state variables in a way that influences the policy, such as by leading the policy to systematically withhold services from an individual, unfairness will result. For example, Hispanics may under-report their pain levels due to cultural factors, misleading a fairness-blind RL agent to assign less therapist time to them (Piette et al., 2023). Deployment of RL algorithms without careful fairness considerations can raise concerns and erode public trust in high-stakes settings.

To formally define and address the fairness problem in sequential decision-making settings, Wang et al. (2025) extended the concept of single-stage counterfactual fairness (CF) in a structural causal framework (Kusner et al., 2018) to the multi-stage setting and proposed a data preprocessing algorithm that ensures CF. A policy is counterfactually fair if, at every time step, the probability of assigning any action does not change had the individual’s sensitive attribute taken a different value, while holding constant other historical exogenous variables and actions. In this light, the data preprocessing algorithm ensures CF by constructing new state variables that are not impacted by the sensitive attribute(s). The reward values in the data are also preprocessed, but the purpose of preprocessing the rewards is to improve the value of the learned optimal policy rather than ensure CF. We refer interested readers to Wang et al. (2025) for more technical details.

The CFRL library implements the data preprocessing algorithm proposed by Wang et al. (2025) and provides a suite of tools to evaluate the value and counterfactual unfairness level achieved by any given policy. Here, “CFRL” stands for “Counterfactual Fairness in Reinforcement Learning”. The library produces preprocessed trajectories that can be used by an off-the-shelf offline RL algorithm, such as fitted Q-iteration (FQI) (Riedmiller, 2005), to learn an optimal CF policy. The library can also simply read in any policy following a required format and return its value (i.e. expected discounted cumulative reward) and counterfactual unfairness level in the environment of interest, where the environment can be either pre-specified or learned from

44 the data.

45 Statement of Need

46 Many existing Python libraries implement algorithms designed to ensure fairness in machine
 47 learning. For example, Fairlearn (Weerts et al., 2023) and aif360 (Bellamy et al., 2018)
 48 provide tools for mitigating bias in single-stage machine learning predictions under statistical
 49 association-based fairness criteria such as demographic parity and equal opportunity. However,
 50 existing libraries do not focus on counterfactual fairness, which defines an individual-level
 51 fairness concept from a causal perspective, and they cannot be easily extended to the generalized
 52 reinforcement learning setting. Scripts available from ml-fairness-gym (D'Amour et al., 2020)
 53 allow users to simulate unfairness in sequential decision-making, but they neither implement
 54 algorithms that reduce unfairness nor address CF. To our knowledge, Wang et al. (2025) is
 55 the first work to study CF in RL. Correspondingly, CFRL is also the first code library to address
 56 CF in the RL setting.

57 The contribution of CFRL is two-fold. First, CFRL implements a data preprocessing algorithm
 58 that ensures CF in offline RL. For each individual in the data, the preprocessing algorithm
 59 sequentially estimates the counterfactual states under different sensitive attribute values and
 60 concatenates all of the individual's counterfactual states at each time point into a new state
 61 vector. The preprocessed data can then be directly used by existing RL algorithms for policy
 62 learning, and the learned policy should be counterfactually fair up to finite-sample estimation
 63 accuracy. Second, CFRL provides a platform for assessing RL policies based on CF. After passing
 64 in any policy and a data trajectory from the environment of interest, users can estimate the
 65 value and counterfactual unfairness level achieved by the policy in the environment of interest.

66 High-level Design

67 The CFRL library is composed of 5 major modules. The functionalities of the modules are
 68 summarized in the table below.

Module	Functionalities
reader	Implements functions that read tabular trajectory data from either a .csv file or a pandas.DataFrame into an array format required by CFRL. Also implements functions that export trajectory data to either a .csv file or a pandas.DataFrame.
preprocessor	Implements the data preprocessing algorithm introduced in Wang et al. (2025).
agents	Implements a fitted Q-iteration (FQI) algorithm (Riedmiller, 2005), which learns RL policies and makes decisions based on the learned policy. Users can also pass a preprocessor to the FQI; in this case, the FQI will be able to take in unprocessed trajectories, internally preprocess the input trajectories, and directly output counterfactually fair policies.
environment	Implements a synthetic environment that produces synthetic data as well as a simulated environment that estimates and simulates the transition dynamics of the unknown environment underlying some real-world RL trajectory data. Also implements functions for sampling trajectories from the synthetic and simulated environments.
evaluation	Implements functions that evaluate the value and counterfactual unfairness level of a policy. Depending on the user's needs, the evaluation can be done either in a synthetic environment or in a simulated environment.

69 A general CFRL workflow is as follows: First, simulate a trajectory using environment or read in

a trajectory using reader. Then, train a preprocessor using preprocessor and preprocess the training trajectory data. After that, pass the preprocessed trajectory into the FQI algorithm in agents to learn a counterfactually fair policy. Finally, use functions in evaluation to evaluate the value and counterfactual unfairness level of the trained policy.

In addition, CFRL also provides tools to check for potential non-convergence that arises in the training of neural networks, FQI, and FQE. More discussions about the sources, checks, and fixes of non-convergence in CFRL can be found in the “Common Issues” section of the CFRL documentation.

Data Example

We provide a data example to demonstrate how CFRL learns a counterfactually fair policy from real-world trajectory data with unknown underlying transition dynamics. We also show how CFRL evaluates the value and counterfactual unfairness level of the learned policy. We note that this is only one of the many workflows that CFRL can perform. For example, CFRL can also generate synthetic trajectory data and use it to evaluate the value and counterfactual unfairness level resulting from some custom data preprocessing methods. We refer interested readers to the “Example Workflows” section of the CFRL documentation for more workflow examples.

We also recorded the computing times of different workflows under different combinations of the number of individuals (N) and the number of transitions (T) in the “Computing Times” section of the CFRL documentation. For example, under $N = 500$ and $T = 10$, the workflow presented in this data example (which is the “real data workflow” in the documentation) ran for 378.6 seconds on average in our computing environment.

Load Data

In this demonstration, we use an offline trajectory generated from a SyntheticEnvironment following some pre-specified transition rules. Although the data is actually synthesized, we treat it as if it is from some unknown environment for pedagogical convenience.

The trajectory contains 500 individuals (i.e. $N = 500$) and 10 transitions (i.e. $T = 10$). The sensitive attribute variable and the state variable are both univariate. The sensitive attribute is binary (0 or 1). The actions are also binary (0 or 1) and were sampled using a behavior policy that selects 0 or 1 randomly with equal probability. The trajectory is stored in a tabular format in a .csv file. We use read_trajectory_from_csv() to load the trajectory from the .csv format into the array format required by CFRL.

```
zs, states, actions, rewards, ids = read_trajectory_from_csv(
    path='../data/sample_data_large_uni.csv', z_labels=['z1'],
    state_labels=['state1'], action_label='action', reward_label='reward',
    id_label='ID', T=10)
```

We then split the trajectory data into a training set (80%) and a testing set (20%) using scikit-learn’s train_test_split(). The training set is used to train the counterfactually fair policy, while the testing set is used to evaluate the value and counterfactual unfairness level achieved by the policy.

```
(zs_train, zs_test, states_train, states_test,
    actions_train, actions_test, rewards_train, rewards_test
) = train_test_split(zs, states, actions, rewards, test_size=0.2)
```

Train Preprocessor & Preprocess Trajectories

We now train a SequentialPreprocessor and preprocess the trajectory. The SequentialPreprocessor ensures the learned policy is counterfactually fair by constructing new state variables that

are not impacted by the sensitive attribute. Due to limited trajectory data, the data to be preprocessed will also be the data used to train the preprocessor, so we set `cross_folds=5` to reduce overfitting. In this case, `train_preprocessor()` will internally divide the training data into 5 folds, and each fold is preprocessed using a model that is trained on the other 4 folds. We initialize the `SequentialPreprocessor`, and `train_preprocessor()` will take care of both preprocessor training and trajectory preprocessing.

```
sp = SequentialPreprocessor(z_space=[[0], [1]], num_actions=2, cross_folds=5,
                           mode='single', reg_model='nn')
states_tilde, rewards_tilde = sp.train_preprocessor(
    zs=zs_train, xs=states_train, actions=actions_train, rewards=rewards_train)
```

Counterfactually Fair Policy Learning

Next, we train a counterfactually fair policy using the preprocessed data and FQI with `sp` as its internal preprocessor. By default, the input data will first be preprocessed by `sp` before being used for policy learning. However, since the training data `state_tilde` and `rewards_tilde` are already preprocessed in our case, we set `preprocess=False` during training so that the input trajectory will not be preprocessed again by the internal preprocessor (i.e. `sp`).

```
agent = FQI(num_actions=2, model_type='nn', preprocessor=sp)
agent.train(zs=zs_train, xs=states_tilde, actions=actions_train,
            rewards=rewards_tilde, max_iter=100, preprocess=False)
```

SimulatedEnvironment Training

Before moving on to the evaluation stage, there is one more step: We need to train a `SimulatedEnvironment` that mimics the transition rules of the true environment that generated the training trajectory, which will be used by the evaluation functions to simulate the true data-generating environment. To do so, we initialize a `SimulatedEnvironment` and train it on the whole trajectory data (i.e. training set and testing set combined).

```
env = SimulatedEnvironment(num_actions=2, state_model_type='nn',
                           reward_model_type='nn')
env.fit(zs=zs, states=states, actions=actions, rewards=rewards)
```

Value and Counterfactual Unfairness Level Evaluation

We now use `evaluate_value_through_fqe()` and `evaluate_fairness_through_model()` to estimate the value and counterfactual unfairness level achieved by the trained policy when interacting with the environment of interest, respectively. The counterfactual unfairness level is represented by a metric from 0 to 1, with 0 representing perfect fairness and 1 indicating complete unfairness. We use the testing set for evaluation.

```
value = evaluate_reward_through_fqe(zs=zs_test, states=states_test,
                                   actions=actions_test, rewards=rewards_test, policy=agent, model_type='nn')
cf_metric = evaluate_fairness_through_model(env=env, zs=zs_test, states=states_test,
                                           actions=actions_test, policy=agent)
```

The estimated value is 7.358 and the CF metric is 0.042, which indicates our policy is close to being perfectly counterfactually fair. Indeed, the CF metric should be exactly 0 if we know the true dynamics of the environment of interest; the reason why it is not exactly 0 here is that we need to estimate the dynamics of the environment of interest during preprocessing, which can introduce finite-sample errors.

Comparisons Against Baseline Methods

We can compare the sequential data preprocessing method in CFRL against a few baselines: Random, which selects each action randomly with equal probability; Full, which uses all

variables, including the sensitive attribute, for policy learning; and Unaware, which uses all variables except the sensitive attribute for policy learning. We implemented these baselines and evaluated their values and counterfactual unfairness levels as part of the code example of the “Assessing Policies Using Real Data” workflow in the “[Example Workflows](#)” section of the CFRL documentation. We summarize below the values and CF metrics calculated in this code example, where “ours” stands for outputs from the SequentialPreprocessor.

	Random	Full	Unaware	Ours
Value	−1.444	8.606	8.588	7.358
Counterfactual Unfairness Level	0	0.407	0.446	0.042

By definition, the “random” baseline always achieves perfect CF. On the other hand, “ours” resulted in much fairer policies than “full” and “unaware”, which suggests that the SequentialPreprocessor can effectively improve CF. Nevertheless, as a trade-off for higher CF, “ours” achieved a lower value than “full” and “unaware”.

Conclusions

CFRL is a Python library that enables CF reinforcement learning through data preprocessing. It also provides tools to calculate the value and unfairness level of a given policy. To our knowledge, it is the first library to address CF problems in the context of RL. The practical utility of CFRL can be further improved via extensions. First, the current CFRL implementation requires every individual in the offline dataset to have the same number of time steps. Extending the library to accommodate variable-length episodes can improve its flexibility and usefulness. Second, CFRL can further combine the preprocessor with popular offline RL algorithm libraries such as d3rlpy (Seno & Imai, 2022), or connect the evaluation functions with established RL environment libraries such as gym (Towers et al., 2024). Third, generalization to non-additive counterfactual states reconstruction can make CFRL theoretically more versatile. We leave these extensions to future updates.

References

- Bellamy, R. K. E., Dey, K., Hind, M., Hoffman, S. C., Houde, S., Kannan, K., Lohia, P., Martino, J., Mehta, S., Mojsilovic, A., Nagar, S., Ramamurthy, K. N., Richards, J., Saha, D., Sattigeri, P., Singh, M., Varshney, K. R., & Zhang, Y. (2018). *AI Fairness 360: An extensible toolkit for detecting, understanding, and mitigating unwanted algorithmic bias*.
- D’Amour, A., Srinivasan, H., Atwood, J., Baljekar, P., Sculley, D., & Halpern, Y. (2020). Fairness is not static: Deeper understanding of long term fairness via simulation studies. *Proceedings of the 2020 Conference on Fairness, Accountability, and Transparency*, 525–534. <https://doi.org/10.1145/3351095.3372878>
- Kusner, M. J., Loftus, J. R., Russell, C., & Silva, R. (2018). *Counterfactual Fairness*. <https://arxiv.org/abs/1703.06856>
- Piette, J. D., Thomas, L., Newman, S., Marinec, N., Krauss, J., Chen, J., Wu, Z., & Bohnert, A. S. B. (2023). An automatically adaptive digital health intervention to decrease opioid-related risk while conserving counselor time: Quantitative analysis of treatment decisions based on artificial intelligence and patient-reported risk measures. *J Med Internet Res*, 25, e44165. <https://doi.org/10.2196/44165>
- Riedmiller, M. (2005). Neural fitted Q iteration – first experiences with a data efficient neural reinforcement learning method. In J. Gama, R. Camacho, P. B. Brazdil, A. M. Jorge, & L.

- 181 Torgo (Eds.), *Machine learning: ECML 2005* (pp. 317–328). Springer Berlin Heidelberg.
182 ISBN: 978-3-540-31692-3
- 183 Seno, T., & Imai, M. (2022). d3rlpy: An offline deep reinforcement learning library. *Journal of*
184 *Machine Learning Research*, 23(315), 1–20.
- 185 Towers, M., Kwiatkowski, A., Terry, J., Balis, J. U., De Cola, G., Deleu, T., Goulão, M.,
186 Kallinteris, A., Krimmel, M., KG, A., & others. (2024). Gymnasium: A standard interface
187 for reinforcement learning environments. *arXiv Preprint arXiv:2407.17032*.
- 188 Wang, J., Shi, C., Piette, J. D., Loftus, J. R., Zeng, D., & Wu, Z. (2025). *Counterfactually fair*
189 *reinforcement learning via sequential data preprocessing*. <https://arxiv.org/abs/2501.06366>
- 190 Weerts, H., Dudík, M., Edgar, R., Jalali, A., Lutz, R., & Madaio, M. (2023). Fairlearn:
191 Assessing and improving fairness of AI systems. In *Journal of Machine Learning Research*
192 (No. 257; Vol. 24, pp. 1–8).

DRAFT